

Sustituyendo Makefiles con Python y Luigi

Javier Torres

Backend Lead @ CARTO



Mario

Luigi


Todo Videos Imágenes Noticias Maps

Todo Imágenes Vídeos Noticias Maps Más

Aproximadamente 678.000.000 resultados (0,60 segundos)

Aproximadamente 156.000.000 resultados (0,52 segundos)

Interés a lo largo del tiempo ?





```
import imageio, numpy, matplotlib.colors

for frame in imageio.mimread('mario.gif'):
    rgb = frame[:, :, 0:3] / 256
    hsv = matplotlib.colors.rgb_to_hsv(rgb)
    hsv[..., 0] = numpy.where(
        numpy.absolute(hsv[..., 0] - 0.5) < 0.48,
        hsv[..., 0],
        (hsv[..., 0] + 0.3) % 1
    )
    rgb = matplotlib.colors.hsv_to_rgb(hsv) * 256
    numpy.copyto(frame[..., 0:3], rgb, casting='unsafe')

durations = [f.meta.duration/1000 for f in frames]
imageio.mimwrite('luigi.gif', frames, duration=durations)
```





De Makefile a Luigi con un ejemplo

Philosophy

Conceptually, Luigi is similar to [GNU Make](#) where you have certain tasks and these tasks in turn may have dependencies on other tasks. There are also some similarities to [Oozie](#) and [Azkaban](#). One major difference is that Luigi is not just built specifically for Hadoop, and it's easy to extend it with other kinds of tasks.


```
hello: hello.c
```

```
    gcc -o $@ $^
```

```
hello: hello.c
    gcc -o $@ $^
```

```
class InputTask(ExternalTask):
    filename = Parameter()

    def output(self):
        return LocalTarget(self.filename)

class Compile(Task):
    def requires(self):
        return InputTask('hello.c')

    def run(self):
        subprocess.run(['gcc', '-o',
                        self.output().path,
                        self.input().path])

    def output(self):
        return LocalTarget('hello')
```

```
hello: hello.c
    gcc -o $@ $^
```

```
class InputTask(ExternalTask):
    filename = Parameter()

    def output(self):
        return LocalTarget(self.filename)

class Compile(Task):
    def requires(self):
        return InputTask('hello.c')

    def run(self):
        subprocess.run(['gcc', '-o',
                        self.output().path,
                        self.input().path])

    def output(self):
        return LocalTarget('hello')
```

```
hello: hello.c
    gcc -o $@ $^
```

```
class InputTask(ExternalTask):
    filename = Parameter()

    def output(self):
        return LocalTarget(self.filename)
```

```
class Compile(Task):
    def requires(self):
        return InputTask('hello.c')

    def run(self):
        subprocess.run(['gcc', '-o',
                        self.output().path,
                        self.input().path])

    def output(self):
        return LocalTarget('hello')
```

```
hello: hello.c
    gcc -o $@ $^
```

```
class InputTask(ExternalTask):
    filename = Parameter()

    def output(self):
        return LocalTarget(self.filename)

class Compile(Task):
    def requires(self):
        return InputTask('hello.c')

    def run(self):
        subprocess.run(['gcc', '-o',
                        self.output().path,
                        self.input().path])
```

```
def output(self):
    return LocalTarget('hello')
```

```
hello: hello.c
```

```
gcc -o $@ $^
```

```
class InputTask(ExternalTask):  
    filename = Parameter()  
  
    def output(self):  
        return LocalTarget(self.filename)
```

```
class Compile(Task):  
    def requires(self):  
        return InputTask('hello.c')
```

```
def run(self):  
    subprocess.run(['gcc', '-o',  
                    self.output().path,  
                    self.input().path])
```

```
def output(self):  
    return LocalTarget('hello')
```

```
%.o: %.c
    gcc -c -o $@ $<

hello: hello.o hello2.o
    gcc -o $@ $^
```

```
%.o: %.c
    gcc -c -o $@ $<

hello: hello.o hello2.o
    gcc -o $@ $^
```

```
class CompileObject(Task):
    obj = Parameter()

    def requires(self):
        return InputTask(self.obj+ '.c')

    def run(self):
        subprocess.run(['gcc', '-c', '-o',
                        self.output().path,
                        self.input().path])

    def output(self):
        return LocalTarget(self.obj+ '.o')
```



```
%.o: %.c
```

```
gcc -c -o $@ $<
```

```
hello: hello.o hello2.o
```

```
gcc -o $@ $^
```

```
class CompileObject(Task):
```

```
    obj = Parameter()
```

```
    def requires(self):
```

```
        return InputTask(self.obj+ '.c')
```

```
    def run(self):
```

```
        subprocess.run(['gcc', '-c', '-o',  
                        self.output().path,  
                        self.input().path])
```

```
    def output(self):
```

```
        return LocalTarget(self.obj+ '.o')
```

```
%.o: %.c  
    gcc -c -o $@ $<
```

```
hello: hello.o hello2.o  
    gcc -o $@ $^
```

```
class Compile(Task):
```

```
    def requires(self):  
        return CompileObject('hello'),  
               CompileObject('hello2')
```

```
    def run(self):  
        subprocess.run(  
            ['gcc', '-o', self.output().path] +  
            [i.path for i in self.input()]
```

```
    def output(self):  
        return LocalTarget('hello')
```

```
%.o: %.c
    gcc -c -o $@ $<

hello: hello.o hello2.o
    gcc -o $@ $^
```

```
class Compile(Task):
    def requires(self):
        return CompileObject('hello'),
            CompileObject('hello2')
```

```
def run(self):
    subprocess.run(
        ['gcc', '-o', self.output().path] +
        [i.path for i in self.input()])
```

```
def output(self):
    return LocalTarget('hello')
```

¿Cuándo usar Luigi?

Extensibilidad

Targets

El objetivo puede ser cualquier cosa, no sólo ficheros.

Tasks

Las tareas son código Python. Mucho más que con comandos de shell.

```
class IssueTarget(luigi.Target):
    def __init__(self, repo, number):
        self.repo = repo
        self.number = number

    def exists(self):
        issue = self.repo.issue_set.where(
            models.Issue.number == self.number
        ).first()
        return issue is not None
```

Extensibilidad

Targets

El objetivo puede ser cualquier cosa, no sólo ficheros.

Tasks

Las tareas son código Python. Mucho más que con comandos de shell.

```
class GetPullRequest(luigi.Task):
    repo = luigi.Parameter()
    number = luigi.IntParameter()

    def run(self):
        user = Github().get_user(self.repo.user)
        repo = user.get_repo(self.repo.name)
        issue = repo.get_issue(self.number)
        models.Issue.create(
            repository=self.repo,
            number=self.number,
            name=issue.title
        )
```

Batteries included

Planificador multi-nodo

Tareas largas y pesadas.

Hadoop / HDFS

Trabajar con grandes volúmenes de datos.

Visualizador web

Estado de ejecución, dependencias

Bases de datos / URIs

Soporte para objetivos en MySQL, MSSQL, Postgres, Redis, FTP, SSHFS...

ETL!

La letra pequeña

Parámetros no serializables (multiprocessing)

```
class GetPullRequest(luigi.Task):
    repo = luigi.Parameter()
    number = luigi.IntParameter()

    def run(self):
        Issue.create(repository=self.repo, number=self.number, name="")

    def output(self):
        return IssueTarget(self.repo, self.number)

class DoStuff(luigi.WrapperTask):
    def requires(self):
        repo = Repository.get(1)
        return [GetPullRequest(repo, i) for i in range(10)]
```

```
$ PYTHONPATH=. luigi --module tasks.test DoStuff --local-scheduler
```

```
$ PYTHONPATH=. luigi --module tasks.test DoStuff --local-scheduler
```

```
INFO: Informed scheduler that task    DoStuff__99914b932b    has status    PENDING
```

```
...
```

```
INFO: Running Worker with 1 processes
```

```
...
```

```
INFO: [pid 12476] Worker Worker(salt=873079442, workers=1, host=archie,  
username=javier, pid=12476) done    DoStuff()
```

```
==== Luigi Execution Summary =====
```

```
Scheduled 11 tasks of which:
```

```
* 11 ran successfully:
```

- 1 DoStuff()
- 10 GetPullRequest(repo=1, number=0...9)

```
This progress looks :) because there were no failed tasks or missing  
dependencies
```

```
==== Luigi Execution Summary =====
```

```
$
```

```
$ PYTHONPATH=. luigi --module tasks.test DoStuff --local-scheduler --workers 4
```

```
$ PYTHONPATH=. luigi --module tasks.test DoStuff --local-scheduler --workers 4

INFO: Informed scheduler that task    DoStuff__99914b932b    has status    PENDING
...
INFO: Running Worker with 4 processes
INFO: [pid 12634] Worker Worker(salt=728263954, workers=4, host=archie,
username=javier, pid=12621) running    GetPullRequest(repo=1, number=1)
ERROR: [pid 12633] Worker Worker(salt=728263954, workers=4, host=archie,
username=javier, pid=12621) failedGetPullRequest(repo=1, number=9)
Traceback (most recent call last):
  File "peewee.py", line 2653, in execute_sql
    cursor.execute(sql, params or ())
psycopg2.DatabaseError: error with status PGRES_TUPLES_OK and no message from
the libpq
```

```
class GetPullRequest(luigi.Task):
    repo_id = luigi.IntParameter()
    number = luigi.IntParameter()

    def run(self):
        Issue.create(repository_id=self.repo_id, number=self.number, name="")

    def output(self):
        return IssueTarget(self.repo_id self.number)

class DoStuff(luigi.WrapperTask):
    def requires(self):
        return [GetPullRequest(1, i) for i in range(10)]
```



```
class GetPullRequest(luigi.Task):
    repo = luigi.Parameter()
    number = luigi.IntParameter()

    @models.db.connection_context()
    def run(self):
        Issue.create(repository=self.repo, number=self.number, name="")

    def output(self):
        return IssueTarget(self.repo, self.number)

class DoStuff(luigi.WrapperTask):
    def requires(self):
        repo = Repository.get(1)
        return [GetPullRequest(repo, i) for i in range(10)]
```

Tareas que no siempre devuelven lo mismo

p.ej: actualizaciones

```
class PullRepository(luigi.Task):  
    repo = luigi.Parameter()  
  
    def run(self):  
        git.Repo(self.repo.path()).remote().pull(progress=CloneProgress(self))  
  
    def output(self):  
        return ???
```

```
class PullRepository(luigi.Task):
    repo = luigi.Parameter()
    executed = False

    def run(self):
        git.Repo(self.repo.path()).remote().pull(progress=CloneProgress(self))
        self.executed = True

    def complete(self):
        return self.executed
```

```
class PullRepository(luigi.Task):
    repo = luigi.Parameter()

    def run(self):
        git.Repo(self.repo.path()).remote().pull(progress=CloneProgress(self))
        os.utime(self.repo.path())

    def output(self):
        return RecentlyUpdatedTarget(self.repo.path())

class RecentlyUpdatedTarget(luigi.Target):
    def __init__(self, path):
        self.path = path

    def exists(self):
        return os.path.getmtime(self.path) + 10 > time.time()
```

```
class PullRepository(luigi.Task):
    repo = luigi.Parameter()
    date = luigi.DateParameter(default=datetime.date.today())

    def run(self):
        git.Repo(self.repo.path()).remote().pull(progress=CloneProgress(self))
        os.utime(self.repo.path(), self.date)

    def output(self):
        return RecentlyUpdatedTarget(self.repo.path(), self.date)

class RecentlyUpdatedTarget(luigi.Target):
    def __init__(self, path, date):
        self.path = path
        self.date = date

    def exists(self):
        return os.path.getmtime(self.path) + 10 > self.date
```

¿ Preguntas ?

¡ Gracias !